



VisionFlow

Coordination Engineering for Federated Human–AI Intelligence

Ecosystem Technical Pitch | May 2026

Dr John O'Hare | DreamLab AI | visionflow.info

5

Federated Substrates

CUDA Kernels

MCP Tools

GPU Speedup

Lines of Rust

OSS Repos

92
185+
55x
120K+
6

EXECUTIVE SUMMARY

VisionFlow is an open-source coordination harness for federated human–AI intelligence. It solves the governance gap that emerges when organisations deploy multiple AI agents. Without shared semantics, cryptographic identity, and human-in-the-loop oversight, agents waste 60–80% of tokens on context rediscovery, duplicate reasoning, and contradictory outputs.

Six federated repositories. 120K+ lines of Rust. 92 CUDA kernels. 185+ MCP tools. Five independent substrates mesh through a single cryptographic identity spine (`did:nostr secp256k1`). Every actor (human, agent, server) authenticates at every layer via the same keypair, enabling federation without centralised infrastructure.

For token and model vendors: VisionFlow is the coordination layer that makes extreme token consumption against hard problems rational. Shared ontology eliminates re-derivation. Provenance chains eliminate re-validation. Judgment Broker eliminates decision spin. Result: 3–5× effective token throughput gain, with every token producing auditable, governed, provenance-tracked reasoning.

COORDINATION FAILURE AT SCALE

AI has moved through a clear progression: LLM → chatbot → reasoning → agent → agentics → external harness. Each stage solves the previous stage's limitation and reveals a new one. VisionFlow occupies the next position: the **coordination harness**, providing shared semantics, federated identity, governance, and cryptographic provenance.

73% of frontline AI adoption happens without management sign-off. Your workforce is already building shadow workflows, stitching together AI agents, automating procurement shortcuts. Your organisation is becoming an agentic mesh whether you plan for it or not.

Four failure modes of uncoordinated AI:

- ✗ **Centralised AI** scales tokens but not trust. One vendor, one failure mode, one billing relationship.
- ✗ **Agent frameworks** scale tasks but not governance. Fast and broken is still broken.
- ✗ **Knowledge tools** scale information but not reasoning. More data doesn't mean better decisions.
- ✗ **Collaboration platforms** scale communication but not coordination. More Slack channels don't solve alignment.

ECOSYSTEM ARCHITECTURE: FIVE SUBSTRATES, ONE IDENTITY

solid-pod-rs | Cryptographic Foundation

- Rust port of JavaScriptSolidServer (JSS), 98% parity. 201 rows tracked: 158 present, 5 partial, 3 missing
- W3C Solid Protocol: LDP, WAC access control, conditional requests, glob GET
- Dual auth: NIP-98 Schnorr (stateless) + Solid-OIDC with DPOP (RFC 9449)
- Full identity provider crate: OIDC authorization-code flow, dynamic client registration, WebAuthn passkeys, NIP-07 Schnorr SSO
- HTTP 402 micropayments: WAC PaymentCondition, per-read satoshi costs, MRC20 token minting, BIP-341 anchoring
- ActivityPub federation: inbox/outbox, HTTP Signature verification, follower fan-out
- Git HTTP backend: smart protocol, NIP-98 bridge auth, WAC-integrated
- Embedded NIP-01 Nostr relay with DID Document publication
- Dual compile: native Tokio + WASM Cloudflare Workers
- 870+ tests, SSRF hardened, dotfile allowlist, path traversal protection

Maturity: v0.4.0-alpha.15. ~20K LOC across 6 crates. Feature-complete, pre-1.0 stabilisation.

Licence: AGPL 3.0. Network-service copyleft (upstream JSS). Commercial licensing available via Melvin Carvalho.

nostr-rust-forum | Governance UI & Relay Kit

- 12 Rust crates: core protocol, 5 Cloudflare Workers, Leptos WASM client, mesh federation, rate limiting, config, search, setup
- Judgment Broker decision surface: 1,015-line governance model
- Agent Control Surface Protocol: kinds 31400–31405 (PanelDefinition, PanelState, ActionRequest, ActionResponse, PanelUpdate, PanelRetired)
- REST governance API: 7 NIP-98-gated endpoints (agent registry, broker cases/decisions, roles)
- Passkey-first auth: WebAuthn PRF derives Nostr keys deterministically. No passwords stored
- 3-zone access: public, members, private (runtime configurable)
- Leptos WASM client: 19 pages, 60+ components, governance dashboard, admin panel
- 15 NIP implementations (NIP-01/07/09/29/33/40/42/44/45/50/52/98 + custom 31400–31405)
- Two-tier pod storage: CF Workers (auto-provisioned) + native solid-pod-rs (git-capable)
- Vector search (NIP-50): in-memory cosine k-NN, RVF binary format


Maturity: v3.0.0-rc7. Phase 1 sprint complete (federated NIP-05, pod-resident keys, data export). Actively deployed.

Licence: Dual MIT / Apache-2.0. Permissive. No copyleft obligation on the forum crate itself.

 **DreamLab Edge | Branded Deployment** | React SPA + Leptos WASM forum on Cloudflare Workers | Operator overlay configuration | Live at dreamlab-ai.com | Consumes all substrates as library crates

JUDGMENT BROKER: HUMAN-IN-THE-LOOP GOVERNANCE

Judgment Broker is an emergent property of agents, humans, a knowledge graph, and a relay mesh coordinating through shared cryptographic identity. No single repository owns it.



Every decision is an immutable Nostr event. In regulated industries (pharma, finance, defence), reconstructing decision provenance after the fact costs orders of magnitude more than generating it by construction. Result: **90% reduction in audit reconstruction cost.**

TOKEN ECONOMICS: WHY TOKENMAXING FOR HARD PROBLEMS IS RATIONAL

Cost of not coordinating. A 50-person team running uncoordinated AI agents wastes **5–8× tokens** on context rediscovery, duplicate reasoning, and contradictory outputs. Each agent session starts cold: no shared ontology, no provenance, no memory of what other agents concluded.

Coordination as token multiplier. Shared ontology means agents stop re-deriving domain vocabulary. Provenance chain means agents stop re-validating conclusions. Judgment Broker means agents stop spinning on decisions they lack authority to make. Result: **3–5× effective throughput gain.**

Hard problems justify deep spend.

- Drug discovery: \$2.6B average cost per approved compound
- Climate modelling: \$50M+ per actionable simulation suite
- Creative production: \$100M+ per franchise

Against these stakes, \$10K–100K on coordinated AI reasoning is a rounding error, provided the coordination harness prevents even one wrong conclusion from propagating.



ROI model:

$$ROI = \frac{V \times \text{coordination_mult} \times \text{governance_div}}{\text{token_cost} \times N}$$

Break-even at **3 agents** on any problem valued above **\$50K**. Beyond that, every additional coordinated agent produces net positive value. Ontology, identity spine, and governance plane are shared infrastructure.

COMPETITIVE LANDSCAPE

Every differentiator sits where no competitor can follow. All competing platforms cluster in the commodity zone.

Platform	Crypto Identity	Data Sovereignty	Governance	Formal Reasoning	Federation	OSS
VisionFlow	✓	✓	✓	✓	✓	✓
Google Spark	✗	✗	~	✗	✗	✗
OpenAI Codex	✗	✗	✗	✗	✗	~
Claude Code	✗	~	~	✗	~	~
Devin	✗	✗	~	✗	✗	✗
CrewAI	✗	~	~	✗	~	✓
AutoGen	✗	~	~	✗	~	✓
LangGraph	✗	~	~	✗	~	✓
Cursor / Windsurf	✗	~	~	✗	✗	✗

✓ Protocol-native ~ Partial/add-on ✗ Absent

Formal Reasoning: Solid Red

Every competitor relies on probabilistic LLM inference. VisionFlow's OWL 2 EL subsumption produces provably correct entailments.

Crypto Identity: Solid Red

Only VisionFlow implements cryptographic agent identity at the protocol level. Every other platform ties identity to someone else's server.

Federation: Solid Red

No competitor supports sovereign nodes trusting each other via cryptographic identity rather than shared infrastructure.

REAL-WORLD SCENARIOS

Climate Modelling Consortium

Three universities, two government agencies, one NGO. Each runs VisionClaw with domain-specific ontologies. OWL 2 EL ensures "sea surface temperature anomaly" means the same thing across all. Data stays in sovereign pods; findings surface through the Judgment Broker.

Pharmaceutical Drug Discovery

Biotech + CRO + regulatory consultancy. Literature mining agents parse 50K papers, chemistry agents run ADMET predictions, compliance agents map to ICH guidelines. Each org's agents operate within their pod boundary. The CRO never sees proprietary target lists.

Creative Production at Scale

12-episode series, five time zones. Production ontology maps episodes to scenes to shots to assets. When a VFX shot depends on an unapproved 3D asset, the OWL 2 constraint propagates through the graph and the GPU physics engine makes the blocked dependency visually obvious.

Real-world validation: DreamLab Creative Hub (50-person team, ~998 graph nodes, daily production) | University of Salford (research partnership, semantic force-directed layout) | THG World Record (250+ concurrent XR users)

SCALING MODEL

👤 Single Operator

One Agentbox, standalone mode. Local SQLite beads, local Solid pod, local events. 90+ skills, 185+ tools, privacy filter active. One API key. One container. Minutes to deploy.

👥 Team (50+)

VisionClaw + forum + Agentbox on a shared relay mesh. Shared ontology, Judgment Broker oversight. Every mutation passes through governance. One GPU host + agent container + CF Workers.

🏢 Federated Enterprise

Multiple sovereign Agentbox instances via Nostr relay mesh. Three transport strata: Tailscale (private WireGuard), Nostr relays (censorship-resistant), Cloudflare Tunnels (edge). Trust via did:nostr, not topology.

Open source. Model-agnostic. Federation-ready.

Every token spent through VisionFlow produces auditable, governed, provenance-tracked reasoning.

github.com/DreamLab-AI

visionflow.info

visionflow@thedreamlab.uk